

OVERVIEW

What is Vexyl TTS?

Vexyl TTS is a voice generation tool that turns written lesson text into natural-sounding spoken audio. It was designed specifically for Indian languages and Indian-accented English — unlike most voice tools, which are built primarily for American or British English.

What makes it different from our current voice system is that it understands **how** to say something, not just **what** to say. You can instruct it to sound calm while explaining something complex, serious when delivering an important fact, or warm and encouraging after a student gets an answer right. Our current tool reads every sentence the same way — like a news ticker.

It runs on our own systems, which means student lesson content is never sent to an outside company. This matters for India's data privacy law.

22	69	0	Free
Indian Languages Supported	Pre-built Voice Options	Cost Per Lesson (API fees)	To Use Commercially

CONTEXT

Why are we looking at this now?

Our current voice system reads lesson text in a flat, robotic tone — the same pace, the same pitch, no matter what the content says. For a student listening to a 25-minute lesson narration, this creates fatigue and reduces how much they absorb.

The CEO has identified that for Indian students to engage deeply, lesson audio needs to feel like a real teacher is in the room — someone who slows down before a key statistic, speaks with authority when something matters, and sounds encouraging when a student is struggling.

The core problem, in simple terms: Our current voice reads "73% of students who skip revision fail their exams" the same way it reads "Welcome to today's lesson." Vexyl TTS would deliver the first sentence slowly, seriously, with a natural pause — because we tell it to.

ASSESSMENT

What works in our favour — and what to watch

✓ ADVANTAGES

- **Sounds like an Indian teacher**

Supports Indian-accented English plus Hindi, Tamil, Telugu, Bengali and 18 more languages. No other tool in this category does this natively.

- **Expressive, not robotic**

We can instruct it to sound calm, serious, warm, or formal — set per paragraph, per lesson topic.

- **No ongoing fees**

Unlike tools that charge per word or per minute of audio, Vexyl runs on our own server. Cost is the server, not the usage.

- **Student data stays private**

Lesson text never leaves our system. Compliant with India's Digital Personal Data Protection Act.

- **Free to use in a paid product**

No royalties, no licensing fees — we can charge students without restrictions.

- **Ready for regional expansion**

When we expand to Hindi or Tamil lessons, the same tool handles it without switching providers.

■ CONCERNS

- **Needs better hardware for full speed**

On standard computers, one full lesson's audio takes 35–60 minutes. Clear fix before launch: upgrade to a higher-performance server.

- **Tone control is per-paragraph**

We set the mood for a whole paragraph, not word-by-word. For lessons, paragraph-level is perfectly sufficient.

- **Cannot copy a specific person's voice**

No way to record a real teacher and clone their exact voice. Uses pre-built voice options only.

- **Newer, smaller community**

A relatively new project. If something breaks, there is less external support than with established tools.

- **We manage it ourselves**

Our team installs it, keeps it running, and applies updates. Not a managed cloud service.

REQUIREMENTS

What does it need — and how do we handle it at demo stage?

Each requirement is listed below with a plain-English explanation and our approach to handling it during testing, before we commit to full production use.

REQUIREMENT	WHY IT MATTERS	HOW WE HANDLE IT FOR TESTING
High-performance computing hardware	Without it, generating audio for one full lesson takes 35–60 minutes — too slow for students.	Use short sentences during testing Confirms voice quality and accent without needing expensive hardware yet.
About 6 GB of storage space	The voice model files are large and must be downloaded once before first use.	One-time download — no recurring cost Downloaded once by our team. Sits on our own server permanently.
Linux or Mac operating system	The automated setup only runs on Linux or Mac, not native Windows.	Use our existing cloud server Our backend already runs Linux. No new machines needed.
Runs as a separate background service	Vexyl runs alongside our main app and is called when audio is needed.	Runs on a fixed address on our server Easy to connect, easy to switch off.
Automatic backup voice if Vexyl is unavailable	If Vexyl goes offline during a lesson, students must not hear silence.	Our current Sarvam voice stays as backup Vexyl is tried first. If it fails, Sarvam takes over automatically.

RECOMMENDED APPROACH

What we do, step by step

1	Set up Vexyl on our existing server — 1 day Install the tool using the automated setup. No new accounts or vendors. Runs on our current server at no extra cost.
2	Test voice quality and Indian accent — 1 day Listen to sample sentences. Confirm the accent sounds natural and that calm, serious, and warm tone instructions make a noticeable difference.
3	Confirm backup voice works correctly — half a day Simulate Vexyl being unavailable. Confirm Sarvam takes over cleanly. No student should ever hear silence.
4	Get sign-off from the team Share audio samples — a flat Sarvam narration versus an expressive Vexyl narration — with CEO and stakeholders. Decide: proceed or adjust.
5	Before real students: upgrade to faster hardware Move Vexyl to a higher-performance cloud server (pay-as-you-go). Full lesson audio drops from 60 minutes to under 10 minutes. This happens before any paying student is onboarded.

OUR RECOMMENDATION

Proceed with Vexyl TTS — testing this week

No other tool in this category supports Indian-accented English natively, runs at zero cost per lesson, keeps student data on our own systems, and is free for commercial products — all four at once.

The only concern (speed on standard hardware) is a known, solvable problem that does not affect our ability to evaluate quality during testing. We recommend the two-to-three day test described above, sharing audio samples with stakeholders, and — if the voice quality is approved — proceeding to the hardware upgrade before student onboarding begins.

Indian English — native	■0 per lesson	Data stays on our servers	Free commercial use	2–3 day test window	No long-term commitment needed
-------------------------	---------------	---------------------------	---------------------	---------------------	--------------------------------